


Virtual Memory

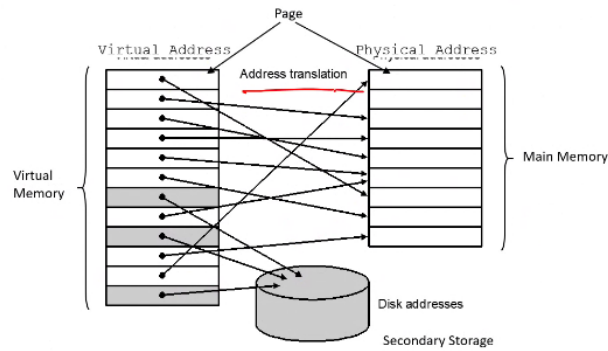
- Motivation: main memory acts as *cache* for secondary storage, e.g., magnetic disk
- Virtual address space, i.e., space addressable by a program is determined by ISA
 - e.g., 64-bit MIPS address space size is 2^{64} – recall jr instruction
 - typically: main memory size \leq disk size \leq virtual address space size
- Program can “pretend” it has main memory of the size of the disk – which is *smaller than the virtual memory* (= whole virtual address space), but *bigger than the actual physical memory* (= DRAM main memory)
 - *Page table* (as we shall see) transparently converts a virtual memory address to a physical memory address, if the data is already in main; if not, it issues call to OS to fetch the data from disk into main
- Virtual memory is organized in fixed-size (power of 2, typically at least 4 KB) blocks, called *pages*. Physical memory is also considered a collection of pages of the same size.
 - the unit of data transfer between disk and physical memory is a page

Demand paging – Bring in the page from SM whenever required!

Logical
to physical




[Block = frame]



Mapping of pages from a virtual address to a physical address or disk address

Page Table Implements Virtual to Physical Address Translation



Page table: page size 4 KB, virtual address space 4 GB, physical memory 1 GB

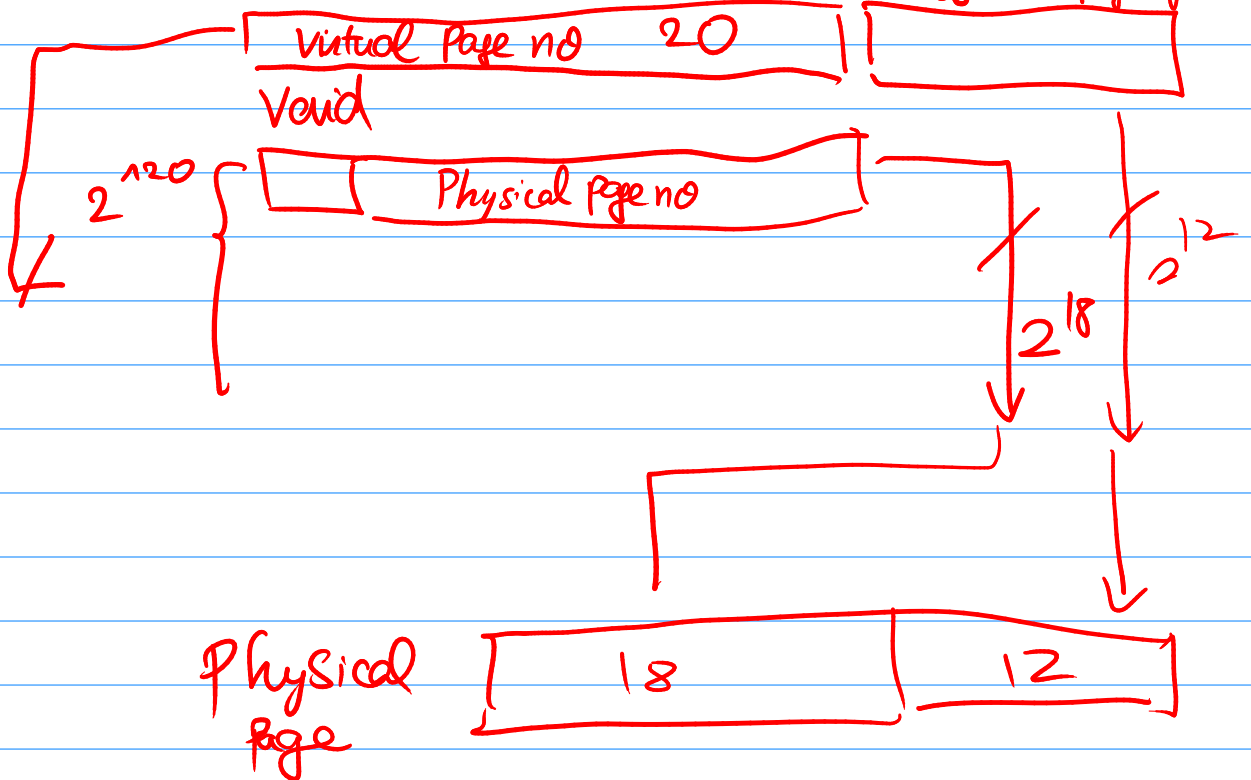
$2^{12} \rightarrow$ page offset

32 bits - VAS

30 bits \rightarrow PAS

\hookrightarrow Physical address space

12 bits page offset



Size = $2^{20} \times (18+1)$

We prefer write back - over write through
Since write back is better in terms of latencies

Write back may cause a problem in redundant disks

TLB

[Page table register]
↓

- ① Common Page table (Too large) ^{requires vast contiguous space}
- ② Separate page table for each process _(lower access time)

32 bit virt mem

4KB Page size $\rightarrow 2^{12}$

4 bytes page table entry

find size

$$2^{20} \times 4$$

= 4 MB page size

4 MB page size is huge

\Rightarrow Each program has its own page table
& the page table register points to the start of the program's
page table (PTR)
= Page table register

- No. of page table entries = address space size / page size

$$= 2^{32} / 2^{12} = 2^{20}$$

- Size of page table = No. of entries \times entry size

$$= 2^{20} \times 4 \text{ bytes} = 4 \text{ MB (huge!)}$$

PTK

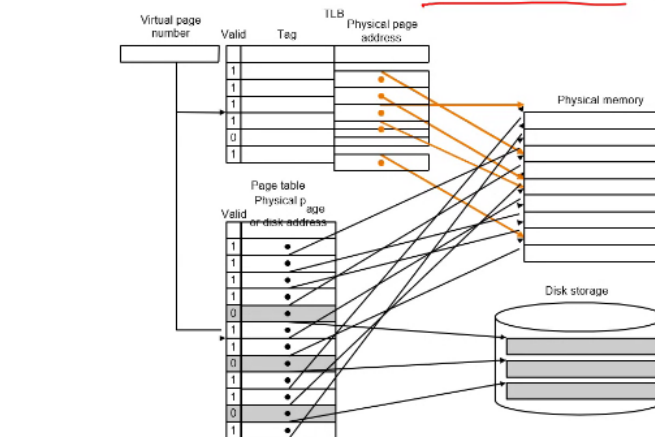
- Note, to avoid large page table size:

- each program has its own page table
 - page table register points to start of program's page table
- to reduce storage required per program page table
 - page table for a program covers the span of virtual memory containing its own code and data
 - other techniques, e.g., multiple-level page tables, hashing virtual address, etc.

Making Address Translation Fast with the Translation-lookaside Buffer

To find page is present

- A cache for address translations – translation-lookaside buffer (TLB):



On a page reference, first look up the virtual page number in the TLB; if there is a TLB miss look up the page table; if miss again then true page fault


$T_{mm} \rightarrow$ memory access
+
 T_{mm} - Page table
2 Times memory access

Page Faults \rightarrow ^{retrieve} not in memory, bring it from disk

Enormous miss penalty $\sim 10^6$ cycles

\Rightarrow 32 or 64 kB page size (large)

All filled



what to replace?

① FCFS (FIFO)

② LRU

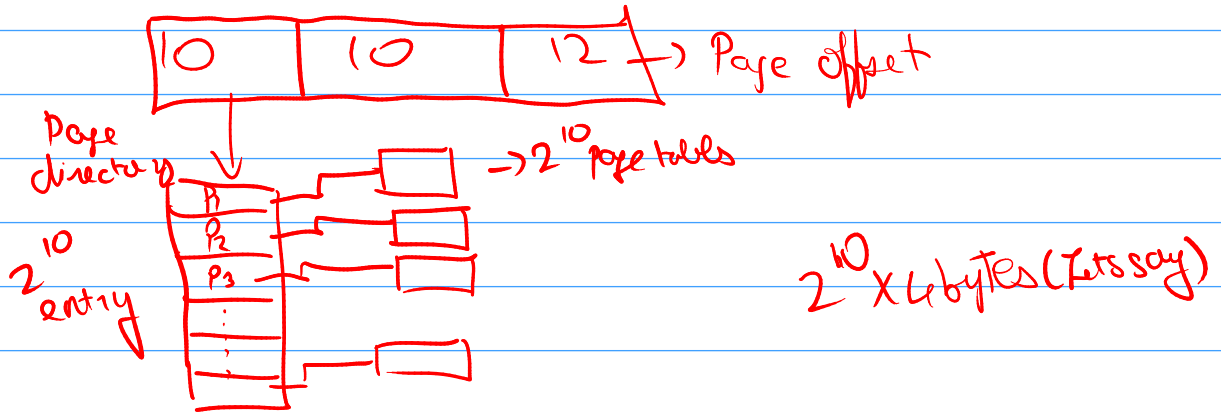
③ Optimal page replacement

(theoretical, best performance in terms of less page misses)

6 is not used anytime soon, it's least used in the future but we cannot see the future

2, 4, 6, 2, 5, 7, 2, 4, 5, 8

↳ replace



we may require all 2^{10} entries in page directory but not all the entries in page tables all at once

→ reducing the no. of page tables in practice

eg: 4 processes will have 4 directory entries

⇒ 4×2^{10} page table entries

if 1 page table entry was 4 bytes long

$$8 \times 2^{10} = 8 \text{ KB}$$

otherwise $2^{20} \times 4 = 4 \text{ MB}$ (too much)

Each process will need its own page table

Page directory is common though, each pointing to a page table of 4kB size.